

IP-SAKTI Sahayak (SIH 2026, PS 26045) — Complete Build & Pitch Playbook

TL;DR

- **What to build:** A multilingual, RAG-based (source-cited) assistant that FIRST classifies an Ayurvedic product into one of six regulatory categories using a deterministic decision tree, THEN answers IP + regulatory questions with an explicit National/International jurisdiction switch, every answer grounded in a cited statute section — never generic LLM chat.
- **The single biggest trap:** TKDL, listed as a data source in the PS, is NOT publicly accessible in 2026 — it remains gated to patent offices under NDA (18 offices as of July 2026), with no public API or subscription live. You must design around this honestly (use open-domain source texts + a "defensive-publication/TKDL-referral" layer) or judges will sink you.
- **How to win:** Scope tightly (2 IP regimes + the classifier + 2 international regimes done deeply beats all 15 regimes done shallow), prove low hallucination with a citation-verification + abstain mechanism, and demo the classifier + cited answer + Hindi/regional language in a 90-second flow.

Key Findings

1. SIH 2026 launched 21 August 2026; internal hackathons Sept 2026, screening Oct 2026, Grand Finale expected December 2026 (36-hour software edition). Judging rewards novelty, complexity, clarity/detail in the prescribed format, feasibility, practicability, sustainability, scale of impact, user experience, and potential for future work.
2. Ministry of Ayush is a repeat SIH sponsor; the SIH 2024 software edition at IIT Tirupati ran an all-AYUSH track (secure startup-registration portal, medicinal-plant/virtual herbal garden, innovation-tracking, annual-report automation). This PS is far more legally complex than any prior Ayush PS.
3. The legal substance is real and fast-moving: Patents Rules 2024, Biological Diversity (Amendment) Act 2023 + Rules 2024 (the AYUSH/codified-TK exemption), the Rule 170 omission saga + Patanjali case, WIPO GRATK Treaty 2024, and EU THMPD 2004/24/EC which effectively blocks Ayurvedic products.
4. The competitive landscape is crowded for general legal AI (Jhana, Lucio, Lexlegis, CaseMine) but EMPTY for an Ayurveda-specific IP+regulatory classifier — that is your moat.
5. Free, hackathon-friendly infrastructure exists end to end: India Code, InPASS, WIPO Lex, CDSCO/FSSAI PDFs, Bhashini/IndicTrans2 for languages, and open-weight LLMs + open vector DBs.

Section 1 — What the Problem Statement Really Asks For

The Ministry's real ask (decomposed): The Ministry of Ayush, through AIIA, wants a trustworthy "first-stop advisor" for the AYUSH community (practitioners, researchers, startups, MSMEs, cultivators) that removes the confusion at the intersection of **IP law + biodiversity/ABS law + drug-regulatory classification**, for both India and export markets. It is NOT asking for a general legal chatbot and NOT asking to reproduce TKDL.

Concretely the deliverables are:

1. **A regulatory classifier first.** Before any IP answer, the tool must ask minimum clarifying questions and classify the formulation into one of six buckets: (a) classical/generic medicine (formula from a First-Schedule authoritative text), (b) patent/proprietary medicine (Section 3(h) D&C Act), (c) new/non-classical drug needing safety+efficacy proof, (d) phytopharmaceutical (Rule 122E / Appendix IB, Schedule Y), (e) Ayurveda-Aahar / nutraceutical (FSSAI), (f) cosmetic. Each bucket has a different IP and ABS posture — e.g., a classical formula hits the Section 3(p) patenting bar and is defended by TKDL, whereas a new drug has real patent potential but must generate clinical evidence.

2. **IP answers with jurisdictional clarity.** Cover 7 IP regimes (patents, GI, trademarks, copyright, designs, trade secrets, plant-variety rights) + ABS duties, and keep National vs International strictly separate via an explicit jurisdiction switch so answers are never conflated.
3. **Source citation always.** Every answer must cite the specific statute/section (not vague summaries) with a link to the authoritative source.
4. **Access facilitation.** Route the user from a question to the right registry/record/form — free official databases directly, and the user's own paid subscriptions ONLY with explicit, logged permission (an audit/consent-log requirement).
5. **Multilingual, plain-language** guidance for the AYUSH community.

What "from the team's side" is required: a working demo of the classifier + cited RAG answer + jurisdiction switch + at least Hindi/one regional language + a consent-logging stub, plus an honest data-sourcing story (especially on TKDL), an accuracy/anti-hallucination story, and a feasible architecture slide. That is the winning minimum.

Section 2 — Non-Negotiable Feature List (multi-angle evaluation)

For each: T=technical, L=legal accuracy, U=user impact, E=evaluator appeal, F=feasibility (1-5, 5=easy).

1. **Deterministic formulation classifier (decision tree/expert system, not pure LLM).**
 - T: A rules engine (YAML/JSON decision tree) + LLM only for parsing user answers. L: highest legal accuracy because category boundaries are defined in statute (First Schedule, Section 3(h), Rule 122E, FSSAI Aahar). U: this is the feature users need most — it resolves the "which regime am I even in?" panic. E: judges love determinism in a legal tool; it is your novelty. F: 5 — trivial to build, huge payoff. **Do not skip.**
2. **Explicit National/International jurisdiction switch with metadata-filtered retrieval.**
 - T: metadata tag on every chunk (jurisdiction, regime, statute, section, date); retrieval filtered by switch. L: prevents the worst error — conflating TRIPS/PCT with Indian Patents Act. U: clarity. E: directly demanded by PS; easy to show. F: 4.
3. **Citation grounding + span-level attribution ("show me the exact section").**
 - T: return the retrieved chunk with statute name + section + URL; highlight the sentence used. L: the core trust mechanism. U: user can verify. E: this is how you beat "why not ChatGPT". F: 4.
4. **Hallucination guardrail with abstain/escalate ("I don't have a cited source for this — consult a professional / here is the registry").**
 - T: if retrieval confidence below threshold or no citation, refuse to answer substantively. L: essential — the Stanford RegLab study showed even Lexis+/Westlaw hallucinate 17-33%. U: safety. E: shows maturity. F: 4.
5. **Multilingual I/O (Hindi + ≥1 regional) via Bhashini/IndicTrans2, translate-query-not-law.**
 - T: translate the user's question to English, retrieve over English legal corpus, answer, translate answer back — but always show the original English statute text to avoid mistranslating law. L: never present a machine-translated statute as authoritative. U: reaches Tier-2/3 founders. E: PS explicitly says multilingual. F: 3.
6. **Access/consent logging for paid-subscription routing.**
 - T: a simple append-only audit log table (user, resource, timestamp, explicit consent yes/no). L: matches PS's "explicit, logged permission". U: trust. E: shows you read the PS carefully. F: 5.

7. Registry/form router (deep links to InPASS, GI Registry, e-Aushadhi, NBA, FSSAI).

- T: a curated link map keyed by classifier output + regime. U: turns advice into action. E: "question to the right form" is in the PS. F: 5.

8. Corpus freshness / update pipeline (dated sources, "as of" stamp on every answer).

- T: each document carries an enactment/amendment date; answers show "as of". L: laws change (Rule 170 omission, 2024 rules). E: pre-empts the "laws go stale" objection. F: 3.

9. ABS / Section 6 & 7 advisor for the biodiversity layer.

- T: rules mapping product type to NBA-approval / SBB-intimation obligations and the codified-TK/AYUSH exemption. L: high-value, directly in PS. F: 4.

10. Disclaimer + "guidance not legal advice" framing on every screen.

- Mitigates unauthorized-practice-of-law risk. F: 5.

Section 3 — Prior Art & Competitive Landscape

General Indian legal AI (crowded, but none Ayurveda-specific): India has more than a thousand legaltech firms tracked in 2026. Leading AI research tools — Jhana.ai (16M+ judgments/statutes, India's "AI paralegal", \$1.6M seed led by Freshworks CEO Girish Mathrubootham), (eCourtsIndia Blog) (jhana) Lucio (founded 2022 in Bengaluru by Vasu Aggarwal & Darsan Guruvayurappan; ~\$7.7M total funding from DeVC and Ashish Kacholia, (eCourtsIndia Blog) including a \$5M seed round in October 2025 — described by Artificial Lawyer as "in effect India's answer to Harvey and Legora"; (Artificial Lawyer) adopted by Trilegal and Cyril Amarchand Mangaldas), (Bar and Bench) Lexlegis.AI (10M+ legal docs, tax focus, US beta "Legal AID" launched at NVIDIA GTC 2026), (eCourtsIndia Blog) (BriefGlance) CaseMine/AMICUS, Manupatra AI, SCC Online — all run RAG over Indian judgments. (iPleaders) **None does Ayurveda IP + regulatory classification.** That gap is your differentiation.

Global legal-AI reliability evidence: The Stanford RegLab / HAI study "Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools" (Magesh, Surani, Dahl, Suzgun, Manning & Ho; arXiv May 2024; published in the Journal of Empirical Legal Studies, 23 April 2025) (LegalAIWorld) found — per its abstract — that "the AI research tools made by LexisNexis (Lexis+ AI) and Thomson Reuters ... each hallucinate between 17% and 33% of the time": (Wiley Online Library) (Stanford) specifically Lexis+ AI 17%, Westlaw AI-Assisted Research 33%, and GPT-4 43%, (AI Law Librarians) despite RAG and "hallucination-free" marketing. Use this to justify your abstain mechanism and citation verification.

WIPO tools (reference/integrate, don't rebuild): WIPO Lex (free IP laws/treaties/judgments, ~48,000 documents, no registration), (Wikipedia) PATENTSCOPE (99M+ patent docs, PCT), (Wipo-analytics) Global Brand Database (75.9M+ records; public API gated to collaborating IP offices), (Patentclient) (Wikipedia) Madrid Monitor, Hague Express, WIPO GRATK Treaty resources.

Government AI stack (reuse): Bhashini / ULCA APIs, AI4Bharat IndicTrans2 (22 scheduled languages, open, MIT-style licensing), (GitHub) (indiaai) IndicBERT, IndicNER — all suitable for the multilingual requirement at near-zero cost.

AYUSH digital initiatives (context + routing targets): NAMASTE portal (morbidity codes/terminologies, feeding WHO ICD-11 TM2), (Lippincott Williams & ...) AMAR (Ayush Manuscripts Advanced Repository), (Ccras) SAHI, Ayush Research Portal (43,000+ studies), (Aimindia) e-Aushadhi (drug licensing), AYUSH Grid, AHMIS. These are for context/routing, not primary IP data.

TKDL — the crucial finding: Despite the August 2022 Cabinet approval of a paid-subscription "widening access" model, TKDL is STILL restricted to patent offices only in 2026. Per CSIR (csir.res.in, 10 July 2026): "With the

signing of the Agreement with IP Australia, eighteen patent offices now have access to the database under Non-Disclosure Agreements (NDAs)" (Council of Scientific ...) — IP Australia signed 9 July 2026 (NE India Broadcast) and Brazil's INPI also concluded a TKDL Access Agreement in 2026. (Council of Scientific ...) The tkdl.res.in homepage states the full database is "available to Patent Offices only under TKDL Access Agreement." (Tkdl) There is no public sign-up, no fee schedule, and no developer API. Current official stats: over 5.2 lakh formulations and practices transcribed (Council of Scientific ...) (~34 million pages) from Ayurveda, Unani, Siddha, Sowa Rigpa and Yoga; and per CSIR, "over 375 patent applications worldwide revoked, rejected, amended, withdrawn or abandoned based on prior-art evidence sourced from the TKDL." (Council of Scientific ...) **Implication:** you cannot ingest TKDL. But the SOURCE TEXTS TKDL is built from (Ayurveda/Siddha/Unani classical texts) are in the open domain and can be sourced independently (csir) — and the First-Schedule authoritative books list is public. Your honest design: (a) reference TKDL as a defensive-publication mechanism your classifier points to; (b) build your own small open-domain corpus of classical-text references; (c) offer a "TKDL access request" referral. Say this out loud in the pitch — it converts your biggest weakness into a credibility win.

Section 4 — Complete Dataset Table (sources, access, licensing, effort)

#	Dataset	URL	Access method	Licensing/legal status	Format	Effort
1	India Code (all central statutes)	indiacode.nic.in	Scrape/download; no formal API	Govt public domain	PDF/HTML	Low-Med
2	Patents Act 1970 + Patents Rules (incl. 2024)	ipindia.gov.in	Download PDFs	Govt open	PDF	Low
3	IP India InPASS (patents)	iprsearch.ipindia.gov.in/publicsearch	Web UI only, CAPTCHA, no official API (Quick Company)	Public data, ToS limits automation	HTML	Med (route, don't bulk-scrape)
4	GI Registry	ipindia.gov.in (GI)	Journal PDFs + registry listing	Public	PDF	Low-Med
5	Trade Marks / Designs registers	ipindia.gov.in	Web UI	Public, ToS limits	HTML	Med
6	Biological Diversity Act 2023 + Rules 2024	nbaindia.org / egazette.gov.in	Download	Govt open	PDF	Low
7	D&C Act 1940 & Rules 1945 (First Schedule, Schedule T, Rule 122E, 158B)	cdsco.gov.in / indiacode	Download	Govt open	PDF	Low
8	DMR (Objectionable Advertisements) Act 1954 + Schedule	indiacode.nic.in	Download	Govt open	PDF	Low
9	FSSAI Ayurveda Aahar Regulations 2022	fssai.gov.in	Download	Govt open	PDF	Low
10	CDSCO phytopharmaceutical guidance/FAQ	cdsco.gov.in	Download	Govt open	PDF/DOC	Low
11	WIPO Lex (TRIPS, CBD, Nagoya, PCT, Madrid, Hague, Budapest, GRATK)	wipo.int/en/web/wipolex	Browse/download; some APIs	Free, no registration	HTML/PDF	Low

#	Dataset	URL	Access method	Licensing/legal status	Format	Effort
12	WIPO GRATK Treaty texts + status	wipo.int (GRATK)	Download	Free	PDF	Low
13	EU THMPD 2004/24/EC + EMA HMPC monographs	eur-lex.europa.eu / ema.europa.eu	Download	EU open	PDF	Low
14	PATENTSCOPE	patentscope.wipo.int	Search + limited bulk ($\leq 10k$), FTP sequences Wipo-analytics	Free	XML/HTML	Med
15	Espacenet/EPO OPS	epo.org	Free API (needs key) Wipo-analytics	Free tier	API	Med
16	Bhashini/ULCA + IndicTrans2	bhashini.gov.in / github AI4Bharat	API + open model weights	Open/MIT-style	API/model	Med
17	AYUSH source classical-text references (open domain)	AMAR ccras.res.in/amar, public texts	Manual curation	Open domain (texts), metadata varies	Text	Med-High
18	TKDL	tkdl.res.in	GATED — patent offices only (18 offices under NDA, July 2026), no public API	Proprietary Govt DB	—	N/A (route/refer only)

Section 5 — Technical Architecture (for the first-round PPT)

One-line architecture: "A rules-first, retrieval-grounded assistant: deterministic classifier → jurisdiction-filtered hybrid RAG over a curated statute corpus → citation-verified, abstaining answer → multilingual output → action router with consent logging."

Data flow (put this diagram in the PPT):

1. **User question** (any Indian language) →
2. **Language layer** (Bhashini/IndicTrans2 detect + translate to English) →
3. **Classifier node** (deterministic decision tree asks minimum clarifying questions; outputs one of 6 categories) →
4. **Jurisdiction switch** (National / International — sets a metadata filter) →
5. **Hybrid retrieval** (BM25 keyword + dense embeddings; section-aware chunking of statutes; metadata filter by jurisdiction+regime+category) →
6. **Reranker** (cross-encoder, e.g., BGE-reranker) →

7. **Grounded generation** (open-weight LLM constrained to answer only from retrieved chunks, must emit statute+section+URL) →
8. **Citation verifier + abstain gate** (if no valid citation → abstain/escalate) →
9. **Answer + "as of" date + disclaimer** → translate back to user's language →
10. **Action router** (deep link to correct registry/form) + **consent log** for any paid-subscription access.

Recommended free/low-cost stack: Vector DB — Chroma or pgvector (Qdrant if you want speed); Embeddings — BGE-M3 or multilingual-e5 (run free/local); Reranker — BGE-reranker; LLM — an open-weight model (e.g., a Llama/Qwen-class model) via a free API or local, with a hosted API as fallback; Orchestration — LangChain/LlamaIndex; Backend — FastAPI; Frontend — Streamlit or React; Eval — RAGAS/TruLens for a small faithfulness scorecard.

Why RAG + rules, not fine-tuning or pure LLM: RAG lets you cite exact sections and update the corpus when laws change (crucial: Rule 170 omission, 2024 rules); the classifier is deterministic because statutory category boundaries are fixed and must not be hallucinated. This hybrid is your feasibility and accuracy story.

Feasibility justification for the PPT: every dataset is free/public; every model has a free tier or open weights; the classifier is a config file; the corpus for a demo is ~20-40 curated statute sections, not the whole legal universe. Buildable by 6 students in the internal-round timeframe.

Section 6 — 36-Hour Grand Finale Execution Plan

Pre-build BEFORE finale (allowed — you refine, not restart):

- Curated corpus (chunked + metadata) for 2 IP regimes (patents + GI) + ABS + the 6-category classifier + 2 international regimes (TRIPS/PCT + EU THMPD).
- Working classifier decision tree, retrieval+rerank pipeline, citation verifier, consent-log schema, base UI, Hindi + one regional language path.
- A seeded evaluation set (~30 Q&A with known correct citations) for a live accuracy scorecard.
- Offline fallback: local embeddings + local LLM + local vector DB so a venue internet failure doesn't kill the demo. Pre-recorded backup demo video and backup laptop.

Build LIVE during 36 hours (to show judges you can absorb feedback):

- Add whichever regime/language/feature the judges request (architecture makes this a config+corpus add, not a rewrite).
- Improve UI polish, add 1-2 more registry routers, expand the eval set live.
- Tune the abstain threshold and show before/after hallucination numbers.

Hour-by-hour (software 36h): H0-2 align on judge feedback + freeze scope; H2-10 integrate the top requested regime/feature; H10-16 multilingual + citation UI polish; H16-22 evaluation scorecard + abstain tuning; H22-28 registry router + consent log demo; H28-32 full dry-runs; H32-36 buffer, backup video, pitch rehearsal.

Minimum demoable slice (the 90-second flow): user types a Hindi question about protecting a home-made herbal formula → classifier asks 2 questions → labels it "classical medicine" → explains Section 3(p) bar + TKDL defensive role, cites the exact section with link, shows "as of" date → user flips to International switch → shows PCT/EU THMPD posture → routes to the right registry → all logged. If nothing else works, this must.

Absorbing judge feedback: explicitly say in round 1 "our architecture is modular — new regimes are added as corpus + metadata + a classifier branch, so any regime you prioritize we can add in hours." That sentence is what makes judges believe the 36-hour claim.

Section 7 — Likely Judge Questions & Prepared Answers

1. **"TKDL isn't public — how will you use it?"** → "Correct, and we verified it: TKDL is still restricted to patent offices under NDA in 2026 (18 offices as of July 2026), [\(Council of Scientific ...\)](#) no public API. We do NOT ingest TKDL. We use the open-domain classical source texts and the public First-Schedule list, and we treat TKDL as a defensive-publication mechanism our tool explains and refers users to." (This honesty wins.)
2. **"Why not just ChatGPT/Gemini?"** → "General LLMs hallucinate law — Stanford found even Lexis+/Westlaw hallucinate 17–33% and GPT-4 43%. [\(AI Law Librarians\)](#) We cite the exact section, verify the citation, and abstain when unsure. Plus we do the Ayurveda-specific classification no general tool does."
3. **"How do you prove accuracy?"** → live RAGAS-style faithfulness scorecard on our 30-question set + citation verification + abstain rate.
4. **"Laws change — stale corpus?"** → every document is dated; answers show "as of"; update pipeline re-ingests changed statutes (we cite the Rule 170 omission and the 2024 rules as live examples).
5. **"Is this legal advice / liability?"** → framed as "guidance, not legal advice," disclaimer on every screen, routes to professionals for filings.
6. **"Scope is enormous — did you under-deliver?"** → "We deliberately built 2 IP regimes + ABS + classifier + 2 international regimes deeply and correctly rather than 15 shallowly. Architecture adds the rest as config."
7. **"Multilingual legal translation risk?"** → we translate the question, not the statute; the authoritative English section is always shown.
8. **"Sustainability after hackathon?"** → runs on free/open infra; AIIA/Ministry can host; corpus maintainable by a small team; consent-logging + audit ready for real deployment.

Section 8 — The Legal Substance You Must Get Right (crib sheet for the corpus)

- **Patents Act 1970 — Section 3(p):** an invention that is "in effect, traditional knowledge or ... aggregation/duplication of known properties of traditionally known component(s)" is NOT patentable [\(Grokikipedia\)](#) — this is the bar a classical formulation hits. Also load 3(e) (mere admixture), 3(i) (methods of treatment), 3(j) (plants/animals).
- **Patents (Amendment) Rules 2024 (in force 15 March 2024):** Request-for-Examination window cut from 48 → 31 months; [\(Obhan Mason\)](#) Form 3 (foreign-filing details) now filed once and updated within 3 months of the First Examination Report instead of repeatedly; [\(Obhan Mason\)](#) Form 27 working statement now filed once every 3 financial years (not annually) and no longer requires revenue/value figures but must state whether the patent is available for licensing; [\(AZB\)](#) relaxed extensions; divisional-application clarification. [\(Anuation\)](#)
- **Biological Diversity (Amendment) Act 2023 + Rules 2024:** Section 6 (IP applications on Indian bio-resources/associated TK need NBA approval; foreign entities need prior approval), [\(Biospectrum India\)](#) Section 7 (Indian entities intimate the State Biodiversity Board before commercial use). [\(Lexology\)](#) The headline relief for this PS: registered AYUSH practitioners, vaid/hakims, users of **codified traditional knowledge** [\(eGazette\)](#) (defined as knowledge in the Ayurvedic books listed in the D&C Act 1940 First Schedule) [\(Mongabay\)](#) and cultivated medicinal plants and their products are **exempt** from Section 7 intimation [\(Lexology\)](#) and from benefit-sharing with local communities; [\(Lexology\)](#) offences decriminalised. [\(Lexology\)](#)
- **D&C Act 1940 & Rules 1945:** definition of ASU drug (Section 3(a)) = medicine manufactured exclusively per the formulae in the First-Schedule authoritative texts; [\(CDSCO\)](#) [\(CDSCO\)](#) "patent or proprietary medicine" (Section 3(h)); Schedule T (GMP), Rule 158B (licensing of ASU categories including classical vs patent/proprietary), Schedule E(1).

- **Phytopharmaceutical (Rule 122E, Appendix IB of Schedule Y, 2015/8th amendment):** a purified/standardised fraction with ≥ 4 defined bioactives, (CliniExperts) regulated as a "new drug" (SRS Journal) by CDSCO (not AYUSH) (PubMed Central) — needs safety/efficacy and confirmatory clinical trials. (PubMed Central) (CliniExperts) This is the bridge category with real patent potential.
- **DMR (Objectionable Advertisements) Act 1954 + Schedule of diseases; the Rule 170 saga:** Rule 170 (D&C Rules) required prior approval of AYUSH ads; (Republic World) the Ministry directed non-enforcement (Aug 2023), then omitted the rule (1 July 2024); (Live Law) the Supreme Court stayed the omission (27 Aug 2024) (Live Law) during the Patanjali/IMA misleading-ads contempt case, then vacated the stay and disposed of the IMA plea (11 Aug 2025). (Republic World) Present this as legally unsettled and evolving.
- **FSSAI Ayurveda Aahar Regulations 2022** for the nutraceutical/food bucket.
- **International layer (keep strictly separate):** TRIPS Art. 27.3(b); CBD + Nagoya Protocol (ABS); the **WIPO GRATK Treaty**, adopted by consensus 24 May 2024 in Geneva (AIPPI) — a mandatory patent-disclosure requirement for inventions based on genetic resources/associated TK (wipo) — which is **not yet in force** (per WIPO, it "will come into force three months after there have been 15 ratifications and accessions"; (WIPO) only about four parties, including Malawi and Uganda, (AIPPI) had ratified as of 2026); PCT, Madrid Protocol, Hague Agreement, Budapest Treaty (micro-organism deposits); and **EU THMPD 2004/24/EC**, whose simplified traditional-use registration was designed for European herbal traditions and, per the European Commission's own 2008 communication, is "not appropriate" for holistic systems like Ayurveda (Eurocam) — the reason Indian products struggle in the EU.
- **Landmark biopiracy cases (for the pitch narrative):** turmeric (US Patent 5,401,504 revoked at USPTO in 1997), (Legacy IAS Academy) neem (EPO patent revoked, final 2005), (Legacy IAS Academy) basmati (RiceTec withdrew claims at USPTO). (Legacy IAS Academy) These built the case for TKDL.

Section 9 — Scoping Recommendation (just right — not over/under)

Do (core, non-negotiable): the 6-category classifier; patents (incl. 3(p), 3(e), 3(i), 3(j), 2024 rules) + GI + ABS (BD Act 2023/Rules 2024, Section 6/7 + AYUSH exemption); 2 international regimes done well (TRIPS/PCT + EU THMPD); citation grounding + abstain; Hindi + 1 regional language; consent log; registry router. **Defer (say "roadmap"):** trademarks/designs/copyright/trade-secrets/plant-variety full depth; Madrid/Hague/Budapest depth; Health Canada/TGA/ASEAN/Japan/China market access; all 22 languages; agentic multi-step workflows; GraphRAG. **Don't overcomplicate:** no knowledge-graph, no fine-tuning, no 15-regime corpus, no fancy agent framework for round 1. A clean rules-first RAG with rock-solid citations and one killer demo flow beats an over-engineered half-working system.

Recommendations (staged)

1. **Now → internal round:** build the classifier + patents/GI/ABS corpus + citation + Hindi + one demo flow; nail the PPT architecture slide and the TKDL honesty slide.
2. **If shortlisted → pre-finale:** add TRIPS/PCT + EU THMPD, the evaluation scorecard, offline fallback, consent log, registry router; rehearse the 90-second flow and backup video.
3. **During 36h:** integrate the judges' top requested regime/feature live; show measurable hallucination reduction.
4. **Benchmarks that change the plan:** if faithfulness on your eval set is below ~90%, tighten the abstain threshold before adding scope; if the multilingual path mistranslates a statute, fall back to showing English law only; if internet is unreliable at the venue, switch to the fully local stack.

Caveats

- Some statistics vary by source/date. AYUSH market: government put the total sector at ~US\$43.4B in 2023 (from US\$2.85B in 2014); (Ayurveda Hub) IBEF puts manufacturing at ~US\$24B and services ~US\$26B in 2024 (>US\$50B combined). (Ayurveda Hub) AYUSH-sector MSMEs grew from 38,216 (Aug 2021) to 53,023 (Jan 2023) per the Udyam portal, (SMERGERS) with 900+ DPIIT-registered AYUSH startups. (Invest India) TKDL figures have risen over time (formulations 3.6→5.2 lakh; patent apps stopped 200→375; offices 14→18). Use the most recent official CSIR/Ministry figures and cite "as of".
- SIH 2026 exact finale dates were not fully fixed at time of writing (expected December 2026); confirm on sih.gov.in.
- The Rule 170 position is legally unsettled — omitted July 2024, stayed Aug 2024, stay vacated / plea disposed 2025; present it as evolving.
- WIPO GRATK ratification counts differ slightly across sources (Malawi and Uganda confirmed; some sources add Albania and Peru); treat the treaty as adopted-but-not-in-force and verify current ratifications on wipo.int.
- This report is strategic guidance, not legal advice; verify every statute section against India Code before putting it in the corpus.